

# Vizualisation

---

## Part 4. Reliability

B9 - Visualisation of Massive Data

M-ALG-103

# Reliability

Dissimilarity

Convergence and the law of large numbers

## Dissimilarity

- ▶ Let us consider a **supervised learning** setup
- ▶ In order to evaluate the quality of our model, we need to measure the **discrepancy** between predictions and observations, noted  $d(\tilde{f}(x), y)$  for one given prediction  $\tilde{f}(x)$  and one given observation  $y$ .
- ▶ And we want to minimize the aggregated discrepancy, summed over all the training samples :

$$\sum_{i=1}^n d(\tilde{f}(x_i), y_i) \quad (1)$$

## Dissimilarity

- ▶ Let us consider a **supervised learning** setup
- ▶ In order to evaluate the quality of our model, we need to measure the **discrepancy** between prediction and observation, noted  $d(\tilde{f}(x), y)$ .
- ▶ And we want to minimize the aggregated discrepancy, summed over all the training samples :

$$\sum_{i=1}^n d(\tilde{f}(x_i), y_i) \quad (2)$$

- But what is  $d(\tilde{f}(x), y)$  ?

## Examples of distances

$x = (x_1, \dots, x_p)$  and  $y = (y_1, \dots, y_p)$  are  $p$ -dimensional **vectors**.

## Examples of distances

$x = (x_1, \dots, x_p)$  and  $y = (y_1, \dots, y_p)$  are  $p$ -dimensional **vectors**.

- ▶  $L_2 : \|x - y\|_2 = \sqrt{\sum_{k=1}^p (x_k - y_k)^2}$  (Euclidian distance, 2-norm distance)

## Examples of distances

$x = (x_1, \dots, x_p)$  and  $y = (y_1, \dots, y_p)$  are  $p$ -dimensional **vectors**.

- ▶  $L_2 : ||x - y||_2 = \sqrt{\sum_{k=1}^p (x_k - y_k)^2}$  (Euclidian distance, 2-norm distance)
- ▶  $L_1 : ||x - y||_1 = \sum_{k=1}^p |x_k - y_k|$  (Manhattan distance, 1-norm distance)

## Examples of distances

$x = (x_1, \dots, x_p)$  and  $y = (y_1, \dots, y_p)$  are  $p$ -dimensional **vectors**.

- ▶  $L_2 : ||x - y||_2 = \sqrt{\sum_{k=1}^P (x_k - y_k)^2}$  (Euclidian distance, 2-norm distance)
- ▶  $L_1 : ||x - y||_1 = \sum_{k=1}^P |x_k - y_k|$  (Manhattan distance, 1-norm distance)
- ▶ weighted  $L_1 : \sum_{k=1}^P w_k |x_k - y_k|$



## Examples of distances

$x = (x_1, \dots, x_p)$  and  $y = (y_1, \dots, y_p)$  are  $p$ -dimensional **vectors**.

- ▶ L2 :  $\|x - y\|_2 = \sqrt{\sum_{k=1}^p (x_k - y_k)^2}$  (Euclidian distance, 2-norm distance)
- ▶ L1 :  $\|x - y\|_1 = \sum_{k=1}^p |x_k - y_k|$  (Manhattan distance, 1-norm distance)
- ▶ weighted  $L_1$  :  $\sum_{k=1}^p w_k |x_k - y_k|$
- ▶  $L_\infty$  :  $\max(x_1, \dots, x_n)$  (infinity norm distance)

# Hamming distance

- ▶  $\#\{x_i \neq y_i\}$  (Hamming distance)

## Hamming distance and edit distance

- ▶  $\#\{x_i \neq y_i\}$  (Hamming distance)
- ▶ linked to **edit distance** : used to quantify how dissimilar two strings are by counting the number of operations needed to transform one into the other (several variants exist)

## General definition of a distance

A **distance** on a set  $E$  is an application  $d : E \times E \rightarrow \mathbb{R}_+$  that must :

## General definition of a distance

A **distance** on a set  $E$  is an application  $d : E \times E \rightarrow \mathbb{R}_+$  that must :

- ▶ be **symetric** :  $\forall x, y, d(x, y) = d(y, x)$

## General definition of a distance

A **distance** on a set  $E$  is an application  $d : E \times E \rightarrow \mathbb{R}_+$  that must :

- ▶ be **symetric** :  $\forall x, y, d(x, y) = d(y, x)$
- ▶ **separate the values** :  $\forall x, y, d(x, y) = 0 \Leftrightarrow x = y$

## General definition of a distance

A **distance** on a set  $E$  is an application  $d : E \times E \rightarrow \mathbb{R}_+$  that must :

- ▶ be **symetric** :  $\forall x, y, d(x, y) = d(y, x)$
- ▶ **separate the values** :  $\forall x, y, d(x, y) = 0 \Leftrightarrow x = y$
- ▶ respect the **triangular inequality**  
 $\forall x, y, z, d(x, y) \leq d(x, z) + d(y, z)$

## General definition of a distance

We could verify that :

- ▶ L2 is a distance
- ▶ Hamming is a distance



Examples of usage of L2 in Machine Learning:

- ▶ kmeans
- ▶ k-nearest neighbors

## Examples when L2 is used

When is L2 used in machine learning ?

- ▶ kmeans
- ▶ k-nearest neighbors
- ▶ agglomerative clustering

## k-nearest neighbors

In a **supervised learning** context :

- ▶ Given training samples that are **labeled**
- ▶ in a **classification** context : for a new input  $x$ , find the  $k$  closest neighbors of  $x$ , and choose a class for  $x$  with the majorities of the classes of the training samples
- ▶ in an **regression** context : for a new input  $x$ , find the  $k$  closest neighbors of  $x$ , and return the average of the outputs of the training samples
- ▶ to find the nearest neighbors, we can use the L2 distance.

## k-nearest neighbors

- ▶ Manhattan (L1) can also be used for kNN.

## Back to L2

- ▶ The euclidean distance can also be used to compute **cost functions** in neural networks



$$C = \frac{(y - \tilde{f}(x))^2}{2} \quad (3)$$

## Examples when Hamming is used

- ▶ compare strings
- ▶ actually we can also use it for k-nearest neighbors

## Cross entropy

- ▶ The **cross-entropy** is another cost function that can be used in neural networks doing classification.
- ▶ A neural network is just a kind of function, here we can just think of it as a function  $\tilde{f}(x)$  of an input  $x$ .

## Cross entropy

- ▶ The **cross-entropy** is another cost function that can be used in neural networks doing classification.
- ▶ A neural network is just a kind of function, here we can just think of it as a function  $\tilde{f}(x) \in \mathbb{R}$  of an input  $x$ .
- ▶ Say we have two possible classes :  $y = 0$  or  $y = 1$ , and a large number of inputs  $x_i$ , each labeled with a class  $y_i$ .
- ▶ Instead of using the **quadratic cost**  $C = \frac{(y - \tilde{f}(x))^2}{2}$ , we can use the **cross entropy**.



## Cross entropy

- Say we have two possible classes :  $y = 0$  or  $y = 1$ , and a large number of inputs  $x_i$ , each labeled with a class  $y_i \in \{0, 1\}$ .  
 $x = (x_1, \dots, x_n)$ ,  $y = (y_1, \dots, y_n)$
- **quadratic cost**

$$C(\tilde{f}(x), y) = \sum_{i=0}^n \frac{(y_i - \tilde{f}(x_i))^2}{2} \quad (4)$$

- **cross entropy**

$$C'(\tilde{f}(x), y) = -\frac{1}{n} \sum_{i=0}^n [y_i \log \tilde{f}(x_i) + (1 - y_i) \log(1 - \tilde{f}(x_i))] \quad (5)$$

...

- ...

...

## Cross entropy

- ▶ Is it a cost ?  $x_i \in \mathbb{R}^p$ ,  $\tilde{f}(x_i) \in [0, 1]$  and  $y_i \in \{0, 1\}$

$$C'(\tilde{f}(x), y) = -\frac{1}{n} \sum_{i=0}^n [y_i \log \tilde{f}(x_i) + (1 - y_i) \log(1 - \tilde{f}(x_i))] \quad (7)$$

- ▶ **positivity** :  $C'(\tilde{f}(x), y) \geq 0$

## Cross entropy

- ▶ Is it a cost ?  $x_i \in \mathbb{R}^n$ ,  $\tilde{f}(x) \in [0, 1]$  and  $y_i \in \{0, 1\}$

$$C'(x, y) = -\frac{1}{n} \sum_{i=0}^n [y_i \log \tilde{f}(x_i) + (1 - y_i) \log(1 - \tilde{f}(x_i))] \quad (8)$$

- ▶ **positivity** :  $C'(\tilde{f}(x), y) \geq 0$
- ▶ **symmetry** :  $C'(\tilde{f}(x), y) = C'(y, \tilde{f}(x))$
- ▶ **separation** :  $C'(\tilde{f}(x), y) = 0 \Leftrightarrow \tilde{f}(x) = y$

## Cross entropy

- ▶ The biggest asset of cross entropy is that it prevents from learning slowdown in neural networks
- ▶ With a quadratic cost, it is possible that the magnitude of the gradient is very small (slowdown)
- ▶ The cross entropy can help fix this issue

## Cross entropy

- ▶ The biggest asset of cross entropy is that it prevents from learning slowdown in neural networks
- ▶ With a quadratic cost, it is possible that the magnitude of the gradient is very small (slowdown)
- ▶ The cross entropy can help fix this issue
- ▶ If you are interested in this it is necessary to understand **backpropagation**, and **gradient descent**.
- ▶ interesting ressource : <http://neuralnetworksanddeeplearning.com/chap3.html>

## Kullbach-Leibler Divergence

- ▶ What if you want measure the discrepancy between **distributions**, instead of **vectors** ?

Expected value (esprance)

- For a discrete random variable  $X$  that takes the values  $x_i$  with probability  $p_i$ :

$$E(X) = \sum_{i=1}^n p_i x_i \quad (9)$$

- For a continuous random variable  $X$  with density  $p(x)$ :

$$E(X) = \int xp(x)dx \quad (10)$$



## Expected value (esprance)

### Exercise 1: Computing an expected value

- ▶ For a discrete random variable  $X$  that takes the values  $x_i$  with probability  $p_i$ :

$$E(X) = \sum_{i=1}^n p_i x_i \quad (11)$$

- ▶ For a continuous random variable  $X$  with density  $p(x)$ :

$$E(X) = \int x p(x) dx \quad (12)$$

Compute the expected value of the dice game.

## Variance

$$\text{var}(X) = E\left((X - E(X))^2\right) \quad (13)$$

## Variance and Covariance

$$\text{var}(X) = E\left((X - E(X))^2\right) \quad (14)$$

$$\text{cov}(X, Y) = E\left((X - E(X))(Y - E(Y))\right) \quad (15)$$

## Kullbach-Leibler Divergence

- ▶ What if you want measure the discrepancy between **distributions**, instead of **vectors** ?
- ▶ For instance, you want to fit a distribution of your choice **model** to **empirical data**.
- ▶ The data points  $(x_1, .., x_n)$  are described by an empirical distribution.

## Kullbach-Leibler Divergence

- ▶ The data points  $(x_1, \dots, x_n)$  are described by an empirical distribution.
- ▶ The **Kullbach-Leibler divergence** is a tool to compare distributions.
- ▶ It is not a distance (it is not symmetric).

## Kullbach-Leibler Divergence



$$\mathcal{D}[p||q] = \mathbb{E}_{\sim p}[\log(\frac{p}{q})] \quad (16)$$



For discrete variables

$$\mathcal{D}[p||q] = \sum_i p(i) \log \frac{p(i)}{q(i)} \quad (17)$$



for continuous variables

$$\mathcal{D}[p||q] = \int_X p(x) \log \frac{p(x)}{q(x)} dx \quad (18)$$

## Exercise 2: Fitting a distribution


- ▶ **cd kl\_divergence**
- ▶ A two dimensional dataset is contained in **empirical\_distribution.csv**. It represents the **age distribution** of some group of people. We want to study this age distribution.
- ▶ load it in **fit\_empirical.py**. We will use the functions provided in the file in order to find the best model, meaning here the model  $M$ , such that  $KL(M||\tilde{P})$  is smallest, with  $\tilde{P}$  the empirical distribution of the data.

## Exercice 2 : Fitting a distribution


- ▶ **First step** : choice of the model
- ▶ Plot the histogram of the data : what model seems to be relevant ?





- ▶ **First step** : choice of the model
- ▶ Plot the histogram of the data : what model seems to be relevant ?



- ▶ We will use **normal laws**. We want to find the normal law that is **the closest to the empirical data**
- ▶ We measure the proximity between the model and the empirical data with the KL divergence.







## Prediction error

- ▶ Given a distance  $d(\tilde{f}(x), y)$  between a prediction and an actual label, we want to measure the **statistical error** of our model
- ▶ This would mean the expected value :

$$\mathbb{E}(d(\tilde{f}(x), y)) \quad (19)$$

## Prediction error

- ▶ Given a distance  $d(\tilde{f}(x), y)$  between a prediction and an actual label, we want to measure the **statistical error** of our model
- ▶ This would mean the expected value :

$$\mathbb{E}(d(\tilde{f}(x), y)) \quad (20)$$

- ▶ Can we compute this ?

## Prediction error

- ▶ Given a distance  $d(\tilde{f}(x), y)$  between a prediction and an actual label, we want to measure the **statistical error** of our model
- ▶ This would mean the expected value :

$$\mathbb{E}(d(\tilde{f}(x), y)) \quad (21)$$

- ▶ Can we compute this ? We can **not** compute it, because we do not know the laws of probabilities of  $x$  and  $y$ .

## Prediction error

We must here understand the difference between :

- ▶ The expected value coming from a certain process with some law probability  $\mathbb{E}$  : we can compute it when we have access to the actual probabilities of the process :

$$\mathbb{E}(d(\tilde{f}(x), y)) \quad (22)$$



## Prediction error

We must here understand the difference between :

- ▶ The expected value coming from a certain process with some law probability  $\mathbb{E}$  : we can compute it when we have access to the actual probabilities of the process :

$$\mathbb{E}(d(\tilde{f}(x), y)) \quad (23)$$

- ▶ The empirical mean that we observe from our datapoints  $(x_1, \dots, x_n)$ :

$$\frac{1}{n} \sum_{i=0}^n d(\tilde{f}(x_i), y_i) \quad (24)$$

## Prediction error

We must here understand the difference between :

- ▶ The expected value coming from a certain process with some law probability  $\mathbb{E}$  : we can compute it when we have access to the actual probabilities of the process :

$$\mathbb{E}(d(\tilde{f}(x), y)) \quad (25)$$

- ▶ The empirical mean that we observe from our datapoints  $(x_1, \dots, x_n)$ :

$$\frac{1}{n} \sum_{i=0}^n d(\tilde{f}(x_i), y_i) \quad (26)$$

- ▶ We have direct access to **only one of these** : which one ?

## Prediction error

We must here understand the difference between :

- ▶ The expected value coming from a certain process with some law probability  $\mathbb{E}$  : we can compute it when we have access to the actual probabilities of the process :

$$\mathbb{E}(d(\tilde{f}(x), y)) \quad (27)$$

- ▶ The empirical mean that we observe from our datapoints  $(x_1, \dots, x_n)$ :

$$\frac{1}{n} \sum_{i=0}^n d(\tilde{f}(x_i), y_i) \quad (28)$$

- ▶ We have direct access to **only one of these** : the empirical mean

# Convergence

- ▶ **However**, under some restrictions, the empirical mean of the error **converges** towards the model error.
- ▶ If the samples are **independent and identically distributed**, and  $d(\tilde{f}(x_i), y_i)$  is integrable.
- ▶ This means that there must be no bias in the samples.

# The end

Questions ?